

Hasan Can Aslan - 60453

Name	Mtu	Network	Address	Ipkts	Ierrs	Opkts	Oerrs	Coll
lo0	16384	<Link#1>	localhost	1042147	0	1042147	0	0
lo0	16384	127	224.0.0.251	1042147	-	1042147	-	-
			all-systems.mcast.net					
lo0	16384	localhost	::1	1042147	-	1042147	-	-
			ff02::fb (refs: 1)					
			ff02::2:ff25:e58e (refs: 1)					
			ff01::1 (refs: 1)					
			ff02::1 (refs: 1)					
			ff02::1:ff00:1 (refs: 1)					
lo0	16384	hasan-macbo	fe80::1:1	1042147	-	1042147	-	-
			ff02::fb (refs: 1)					
			ff02::2:ff25:e58e (refs: 1)					
			ff01::1 (refs: 1)					
			ff02::1 (refs: 1)					
			ff02::1:ff00:1 (refs: 1)					
gif0*	1280	<Link#2>		0	0	0	0	0
stf0*	1280	<Link#3>		0	0	0	0	0
api*	1500	<Link#5>	f2:18:98:62:13:01	0	0	0	0	0
en0	1500	<Link#6>	f0:18:98:62:13:01	14531107	0	20312842	75214	0
en0	1500	hasan-macbo	fe80:6:1851:25c0:	14531107	-	20312842	-	-
			ff02::fb (refs: 1)					
			ff01::1 (refs: 1)					
			ff02::2:ffa8a:dd10 (refs: 1)					
			ff02::1 (refs: 1)					
			ff02::1:fff3a:78a4 (refs: 1)					
en0	1500	192.168.1	192.168.1.13	14531107	-	20312842	-	-
			239.255.255.250					
			224.0.0.251					
			all-systems.mcast.net					
en1	1500	<Link#7>	82:dd:b9:a2:4c:01	0	0	0	0	0
en2	1500	<Link#8>	82:dd:b9:a2:4c:00	0	0	0	0	0
en3	1500	<Link#9>	82:dd:b9:a2:4c:05	0	0	0	0	0
en4	1500	<Link#10>	82:dd:b9:a2:4c:04	0	0	0	0	0
bridg	1500	<Link#11>	82:dd:b9:a2:4c:01	0	0	0	0	0
awd10	1500	<Link#12>	1a:43:3d:51:25:33	14761	0	144226	0	0
awd10	1500	fe80::1843:	fe80:c::1843:3dff	14761	-	144226	-	-
			ff01::1 (refs: 1)					
			ff02::1 (refs: 1)					
			ff02::1:fff51:2533 (refs: 1)					
			ff02::2:ffa8a:dd10 (refs: 1)					
			ff02::fb (refs: 1)					
llw0	1500	<Link#13>	1a:43:3d:51:25:33	0	0	0	0	0
llw0	1500	fe80::1843:	fe80:d::1843:3dff	0	-	0	-	-
			ff01::1 (refs: 1)					
			ff02::1 (refs: 1)					
			ff02::1:fff51:2533 (refs: 1)					
			ff02::fb (refs: 1)					
			ff02::1:ffd7:2354 (refs: 3)					
			ff02::1:ff2f:6d75 (refs: 4)					
			ff02::2:ffa8a:dd10 (refs: 11)					
			ff02::1:ffbd:2c9a (refs: 3)					
utun0	1380	<Link#14>		0	0	20	0	0
utun0	1380	hasan-macbo	fe80::e:e457:d006	0	-	20	-	-
			ff01::1 (refs: 1)					
			ff02::2:ffa8a:dd10 (refs: 1)					
			ff02::1 (refs: 1)					
			ff02::1:ffda:fcd79 (refs: 1)					
utun1	2000	<Link#15>		0	0	20	0	0
utun1	2000	hasan-macbo	fe80::f:60e1:e539	0	-	20	-	-
			ff01::1 (refs: 1)					
			ff02::2:ffa8a:dd10 (refs: 1)					
			ff02::1 (refs: 1)					
			ff02::1:fff69:f22e (refs: 1)					
utun2</								

2. Destination column: Shows the destination network.

Gateway column: Indicates the router through which packets are forwarded.

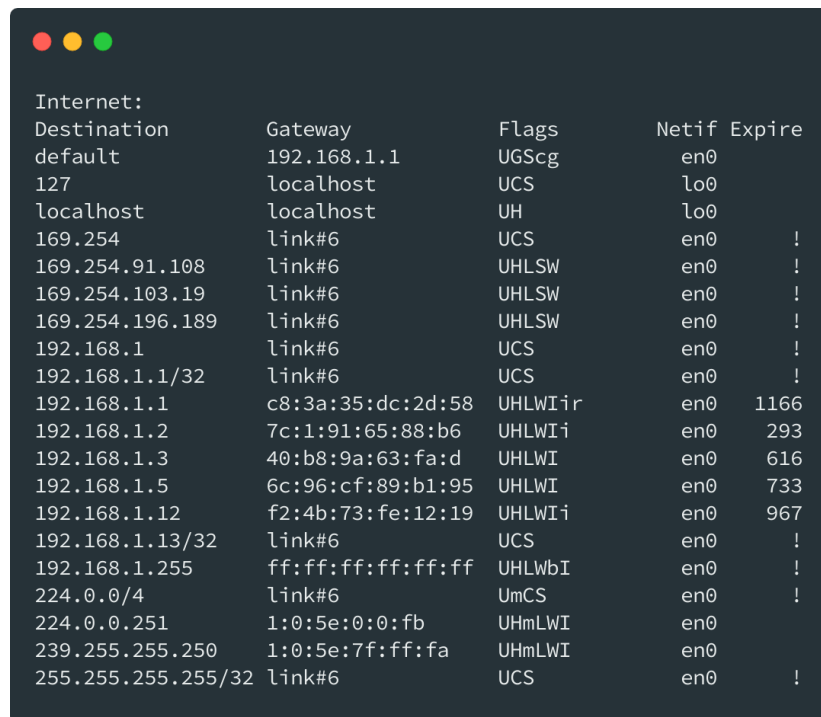
Flags column: The U flag indicates that the route is up route is valid.

The G flag indicates that the route is to a gateway.

The S flag indicates route added with the route command.

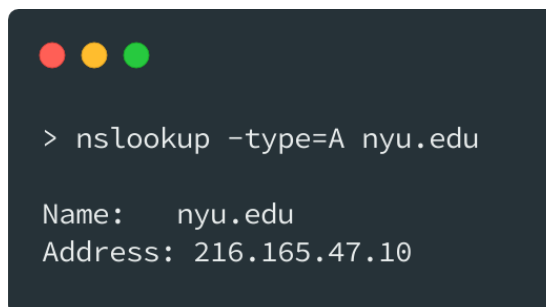
The H flag indicates route is to a host rather than to a network, where the destination address is a complete address.

Netif column: Indicates the interface on the local host that is the source endpoint of the transmission.



Internet:	Destination	Gateway	Flags	Netif	Expire
	default	192.168.1.1	UGScg	en0	
	127	localhost	UCS	lo0	
	localhost	localhost	UH	lo0	
	169.254	link#6	UCS	en0	!
	169.254.91.108	link#6	UHLSW	en0	!
	169.254.103.19	link#6	UHLSW	en0	!
	169.254.196.189	link#6	UHLSW	en0	!
	192.168.1	link#6	UCS	en0	!
	192.168.1.1/32	link#6	UCS	en0	!
	192.168.1.1	c8:3a:35:dc:2d:58	UHLWIir	en0	1166
	192.168.1.2	7c:1:91:65:88:b6	UHLWIi	en0	293
	192.168.1.3	40:b8:9a:63:fa:d	UHLWI	en0	616
	192.168.1.5	6c:96:cf:89:b1:95	UHLWI	en0	733
	192.168.1.12	f2:4b:73:fe:12:19	UHLWIi	en0	967
	192.168.1.13/32	link#6	UCS	en0	!
	192.168.1.255	ff:ff:ff:ff:ff:ff	UHLWbI	en0	!
	224.0.0/4	link#6	UmCS	en0	!
	224.0.0.251	1:0:5e:0:0:fb	UHmLWI	en0	
	239.255.255.250	1:0:5e:7f:ff:fa	UHmLWI	en0	
	255.255.255.255/32	link#6	UCS	en0	!

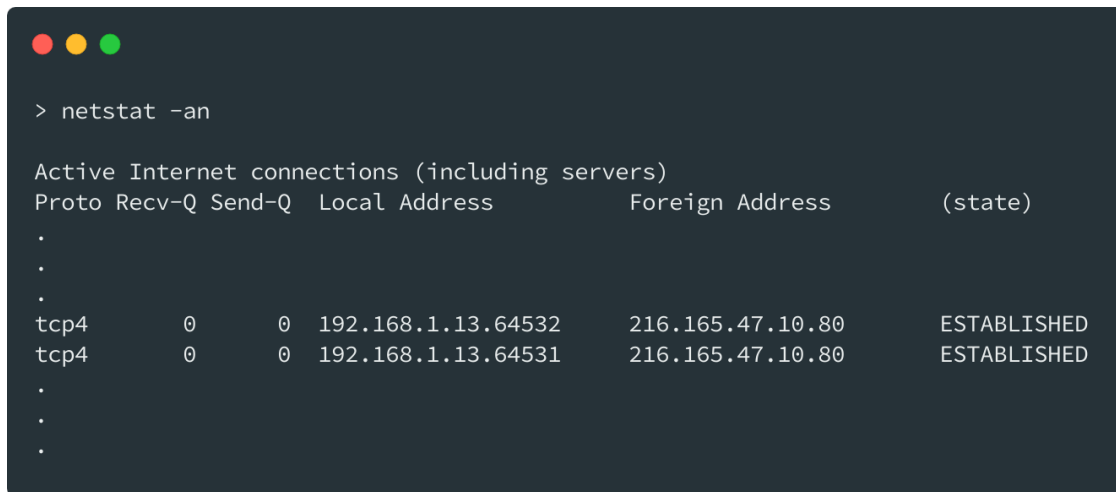
3. My client using ports 64532 and 64531 to connect nyu.edu. First, I get the IP address of the host, using **nslookup**.



```
> nslookup -type=A nyu.edu

Name:    nyu.edu
Address: 216.165.47.10
```

Then using `netstat -P tcp` command, we can find nyu.edu's IP address in the list.

A terminal window with a dark background and three colored window control buttons (red, yellow, green) in the top-left corner. The terminal displays the output of the command `netstat -an`. The output shows active Internet connections, including two established TCP connections to 216.165.47.10.80.

```
> netstat -an

Active Internet connections (including servers)
Proto Recv-Q Send-Q Local Address           Foreign Address         (state)
.
.
.
tcp4      0      0 192.168.1.13.64532      216.165.47.10.80      ESTABLISHED
tcp4      0      0 192.168.1.13.64531      216.165.47.10.80      ESTABLISHED
.
.
.
```

Part 1B. ICMP Analysis

4. TTL, time to live is the number of hops or amount of time that a packet exists inside a network. When TTL exceeded, packet is discarded. Time to live is important for caching information. It determines how long the data can be cached and when the information should be updated. We can find TTL under the network layer, for example under the Internet Protocol.

5. ICMP protocol is designed to communicate between network layers. However, port numbers are used to communicate between application layers, so ICMP has neither destination nor source port number.

6. Minimum TTL is 1.

```
> Frame 126: 182 bytes on wire (1456 bits), 182 bytes captured (1456 bits) on interface en0, id 0
> Ethernet II, Src: TendaTec_dc:2d:58 (c8:3a:35:dc:2d:58), Dst: Apple_62:13:01 (f0:18:98:62:13:01)
> Internet Protocol Version 4, Src: 64.125.26.172, Dst: 192.168.1.13
~ Internet Control Message Protocol
  Type: 11 (Time-to-live exceeded)
  Code: 0 (Time to live exceeded in transit)
  Checksum: 0xbe96 [correct]
  [Checksum Status: Good]
  Unused: 00000000
~ Internet Protocol Version 4, Src: 192.168.1.13, Dst: 216.165.47.10
  0100 .... = Version: 4
  .... 0101 = Header Length: 20 bytes (5)
  > Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
  Total Length: 52
  Identification: 0x9507 (38151)
  > Flags: 0x00
  Fragment Offset: 0
  > Time to Live: 1
  Protocol: UDP (17)
  Header Checksum: 0x5b4d [validation disabled]
  [Header checksum status: Unverified]
  Source Address: 192.168.1.13
  Destination Address: 216.165.47.10
  > User Datagram Protocol, Src Port: 38120, Dst Port: 33465
  > ICMP Multi-Part Extensions
  > Data (24 bytes)
```

7. My computer sends first a packet with $TTL = 1$, the second packet has $TTL = 2$ so and so on. A router will decrement the packet's TTL when the packet passes through. If the packet arrives the router with TTL value equals to 1, then router sends an ICMP packet to the source. In this way we can count the router along the path to the destination.

8. For each router along the path, **traceroute** sends 3 probes.

```
traceroute to nyu.edu (216.165.47.10), 64 hops max, 52 byte packets
 1 192.168.1.1 (192.168.1.1)  9.176 ms  2.175 ms  3.090 ms
 2 212.156.201.180.static.turktelekom.com.tr (212.156.201.180)  21.905 ms  19.413 ms
19.555 ms
 3 81.212.71.241.static.turktelekom.com.tr (81.212.71.241)  13.059 ms  13.178 ms  13.232
ms
 4 01-adana-xrs-t2-1---31-hatay-t3-1.statik.turktelekom.com.tr (195.175.167.56)  19.527
ms  17.302 ms  17.085 ms
 5 35-izmir-xrs-t2-1---01-adana-xrs-t2-1.statik.turktelekom.com.tr (195.175.166.0)
28.367 ms  30.671 ms  31.680 ms
 6 * 35-ebgp-izmir-sr12e-k---35-izmir-xrs-t2-1.statik.turktelekom.com.tr (81.212.30.5)
31.320 ms  29.615 ms
 7 301-fra-col-1---06-ebgp-ulus-sr12e-k.statik.turktelekom.com.tr (212.156.101.32)
73.037 ms  69.714 ms  68.996 ms
 8 62.157.248.1 (62.157.248.1)  868.168 ms * *
 9 pd900cb02.dip0.t-ipconnect.de (217.0.203.2)  85.037 ms  79.513 ms *
10 new-york-un.ear3.newyork1.level3.net (4.28.130.118)  160.253 ms
80.150.170.214 (80.150.170.214)  91.782 ms
new-york-un.ear3.newyork1.level3.net (4.28.130.118)  153.203 ms
11 ae12.cs1.fra6.de.eth.zayo.com (64.125.26.172)  150.142 ms  146.613 ms
dmzgw-p2p-extgwc.net.nyu.edu (128.122.254.75)  160.302 ms
12 nyugwa-ptp-dmzgw-vl3082.net.nyu.edu (128.122.254.110)  157.920 ms
ae2.cs1.ams17.nl.eth.zayo.com (64.125.29.59)  167.329 ms  167.260 ms
13 nyufw-outside-ngfw-vl3080.net.nyu.edu (128.122.254.116)  157.550 ms  160.002 ms *
14 * * *
15 * * wsqdcgwa-vl902.net.nyu.edu (128.122.1.38)  164.298 ms
16 ae5.cs3.lga5.us.eth.zayo.com (64.125.29.126)  148.268 ms * *
17 * * *
18 * 209.66.118.177.idia-282827-zyo.zip.zayo.com (209.66.118.177)  173.936 ms  167.051 ms
19 * * *
20 * * *
21 nyufw-outside-ngfw-vl3080.net.nyu.edu (128.122.254.116)  163.018 ms *  154.461 ms
22 * * *
23 wsqdcgwa-vl902.net.nyu.edu (128.122.1.38)  157.344 ms * *
```

9. I reach to the destination much faster, at 11 hops, when tracing 18.31.0.200.

10. Routing Blackhole is the routing the all network traffic into blackhole and lost it. We can use such a system to prevent attacks to our networks like DDoS attack. However, normal traffic is also affected such prevention system

11. We can arrange the output fields by using the -o flag. I use following command:

```
sudo mtr -o "LSDR NBAW JMXI" nyu.edu
```

Host	Packets				Pings							
	Loss%	Snt	Drop	Rcv	Last	Best	Avg	Wrst	Jttr	Javg	Jmax	Jint
1. 192.168.1.1	3.4%	89	3	86	12.5	2.0	23.7	317.8	7.4	17.6	295.	156.
2. 212.156.201.180.static.turktelekom.com.tr	1.1%	89	1	87	19.6	15.2	50.2	355.5	21.2	23.7	183.	307.
3. 81.212.71.241.static.turktelekom.com.tr	0.0%	88	0	88	13.3	13.1	37.5	301.5	4.6	17.7	144.	188.
4. 01-adana-xrs-t2-1---31-hatay-t3-1.statik.turktelekom.com.tr	0.0%	88	0	88	18.4	15.3	37.5	276.3	7.6	18.7	164.	162.
5. 35-izmir-xrs-t2-1---01-adana-xrs-t2-1.statik.turktelekom.com.tr	0.0%	88	0	88	33.0	29.7	49.1	253.2	6.4	18.2	183.	129.
6. 35-ebgp-izmir-sr12e-k---35-izmir-xrs-t2-1.statik.turktelekom.com.tr	2.3%	88	2	86	35.9	28.9	46.9	210.7	16.6	16.4	167.	218.
7. 301-fra-col-1---06-ebgp-ulus-sr12e-k.statik.turktelekom.com.tr	1.1%	88	1	87	73.4	68.9	85.8	403.2	2.9	20.0	330.	179.
8. 62.67.19.245	96.6%	88	84	3	80.8	77.3	79.2	80.8	3.5	1.8	3.5	5.4
9. ae-1-3501.ear3.newyork1.level3.net	84.1%	88	74	14	164.3	156.6	176.4	354.0	6.5	33.1	197.	252.
10. new-york-un.ear3.newyork1.level3.net	1.1%	88	1	87	160.1	153.3	191.9	549.2	22.5	25.2	302.	245.
11. dmzgw-b-p2p-extgwc.net.nyu.edu	81.6%	88	71	16	157.9	157.0	215.1	479.1	22.9	52.3	322.	443.
12. nyugwa-ptp-dmzgw-b-vl3082.net.nyu.edu	1.1%	88	1	87	159.5	155.5	194.7	487.4	1.1	34.8	250.	460.
13. nyufw-outside-ngfw-vl3080.net.nyu.edu	79.3%	88	69	18	158.3	155.7	169.2	208.5	50.2	17.3	50.2	208.
14. (waiting for reply)												
15. wsqdcgwa-vl902.net.nyu.edu	95.4%	88	83	4	185.2	159.3	271.1	468.2	86.6	148.	308.	542.
16. (waiting for reply)												

These fields are described in man page of `mtr` as follows:

```
+-----+
|L | Loss ratio      |
+-----+
|D | Dropped packets |
+-----+
|R | Received packets |
+-----+
|S | Sent Packets    |
+-----+
|N | Newest RTT(ms)  |
+-----+
|B | Min/Best RTT(ms)|
+-----+
|A | Average RTT(ms) |
+-----+
|W | Max/Worst RTT(ms)|
+-----+
|V | Standard Deviation |
+-----+
|G | Geometric Mean   |
+-----+
|J | Current Jitter   |
+-----+
|M | Jitter Mean/Avg.  |
+-----+
|X | Worst Jitter     |
+-----+
|I | Interarrival Jitter |
+-----+
```

12. **mtr** continuously updates times by polling a remote server and allows us to see how the latency and performance changes over time. On the other hand, **traceroute** allows us to discover the pathway to a host. When we compare two Wireshark record, we can see that **mtr** continuously sends probes whilst **traceroute** sends predetermined number of probes.

Part 2. DV Routing Simulator

Part 2A. Implementation

When constructing Node class, I initialize id from static counter, get cost and neighbors from static arrays of **DVSimulator** class. **myDV** is initially equals to cost of that node. I initialize **bestPath** array based on following conditions: If nodes id or neighbor's id is equal to destination id, then use that id, otherwise get a random node from **randomNeighbor** method.

```
public Node() {
    this.id = count++;
    this.cost = DVSimulator.cost[this.id];
    this.neighbors = DVSimulator.neighbors[this.id];
    this.myDV = this.cost;

    for (int i = 0; i < DVSimulator.NUMNODES; i++) {
        if (i == this.id) {
            this.bestPath[i] = this.id;
        } else if (Node.contains(neighbors, i)) {
            this.bestPath[i] = i;
        } else {
            this.bestPath[i] = randomNeighbor();
        }
    }

    notifyNeighbors();
}
```

For the `notifyNeighbors` method, I create a new packet for each neighbor that created with current node `id` as source, neighbor id as destination and current node's `myDV` array as the dv. Then, I send the packet.

```
public void notifyNeighbors() {
    for (int neighbor : neighbors) {
        Packet packet = new Packet(this.id, neighbor, this.myDV);
        DVSimulator.sendPacket(packet);
    }
}
```

Note: I get the following method from <https://stackoverflow.com/a/34541755>. It is simple method for ease of use when searching a value in array.

```
public static boolean contains(final int[] arr, final int key) {
    return Arrays.stream(arr).anyMatch(i -> i == key);
}
```


Part 2B. Optional (Bonus)

I compare for each value in the DV received from neighbor as **new_cost** and current DV as **current_cost**. If there is a cheaper path, I update **myDV** with corresponding cost and update **bestPath** with corresponding neighbor's id. If there is such a change in current Node's DV I update **hasDVChanged** value as true. If **hasDVChanged** value is true, then I **notifyNeighbors** and increment **numUpdates**. By doing that, I increment and notify neighbors just once for each packet.

```
public void updateDV(Packet p) {
    int neighbor_id = p.getSource();
    neighborDV[neighbor_id] = p.getDV();
    boolean hasDVChanged = false;

    for (int i = 0; i < myDV.length; i++) {
        int new_cost = cost[neighbor_id] + neighborDV[neighbor_id][i];
        int current_cost = myDV[i];

        if (new_cost < current_cost) {
            myDV[i] = new_cost;
            bestPath[i] = neighbor_id;
            hasDVChanged = true;
        }
    }

    if (hasDVChanged) {
        notifyNeighbors();
        numUpdates++;
    }
}
```